

Homework 8 - Predictive Modeling in Finance and Insurance

Dennis Goldenberg

2024-03-31

```
library(readxl)
library(ggplot2)
library(patchwork)
```

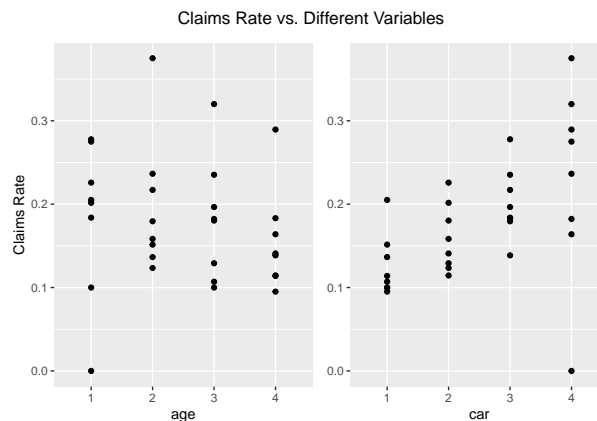
1. Poisson Regression

```
claimsData <- read_excel("Table 9.13 Car insurance.xls", skip = 2,
                        sheet = "Sheet1", .name_repair = "unique_quiet")
claimsData$age <- factor(claimsData$age)
claimsData$district <- factor(claimsData$district)
claimsData$car <- factor(claimsData$car)
claimsData$y <- as.integer(claimsData$y)
```

a. Exploratory Data Analysis

I first plot the claims rate by age, and by car

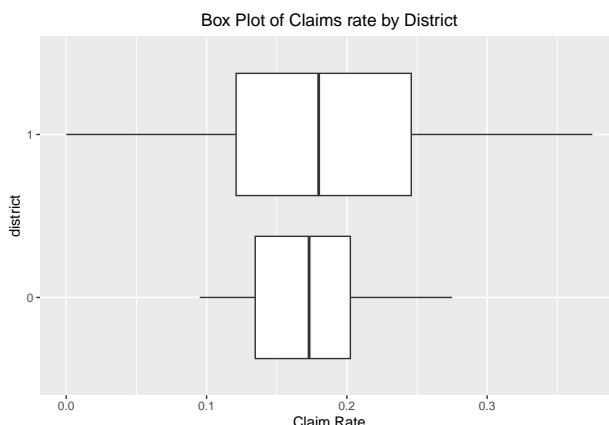
```
claimsData$claimRate <- claimsData$y/claimsData$n
p1 <- ggplot(data = claimsData) + geom_point(aes(age, claimRate)) +
  ylab("Claims Rate")
p2 <- ggplot(data = claimsData) + geom_point(aes(car, claimRate)) +
  theme(axis.title.y = element_blank())
p1 + p2 + plot_annotation(title = "Claims Rate vs. Different Variables",
                        theme = theme(plot.title = element_text(hjust = .5)))
```



Age does not seem to have a very significant impact on claims rate; if anything, there may be a potential slight negative correlation. However, the car variable has a strong positive correlation with claims rate; as

the number of the car goes up, there is a noticeable shift in distribution of claims rates at each level. Next, I compare the boxplot of claims rate by district:

```
ggplot(data = claimsData) +
  geom_boxplot(aes(x = claimRate, y = district, group = district)) +
  labs(title = "Box Plot of Claims rate by District") +
  xlab("Claim Rate") +
  theme(plot.title = element_text(hjust = 0.5))
```



Note that district 0 and district 1 have about the same mean, but district 1's IQR is far larger (as is its overall range), suggesting that claims rates in district 1 have higher variance.

b. Testing for significance of Poisson Regression

I fit both models, and get their log likelihoods:

```
noInter <- glm(y ~ age + car + district + offset(log(n)),
  family = poisson(link = 'log'),
  data = claimsData)
Inter <- glm(y ~ age + car + district + age*car + age*district +
  car*district + offset(log(n)),
  family = poisson(link = 'log'),
  data = claimsData)
ll_noI <- as.numeric(logLik(noInter))
ll_I <- as.numeric(logLik(Inter))
sprintf("Log-Like - No Interactions: %.3f; Interactions: %.3f", ll_noI, ll_I)
```

```
## [1] "Log-Like - No Interactions: -96.035; Interactions: -86.828"
```

Then, I note that $C = 2 \left[\ell(\hat{\beta}_{\text{full}}) - \ell(\hat{\beta}_{\text{reduced}}) \right] \sim \chi^2(q)$, where q is the number of extra parameters in the model (in this case, the number of interactions is 15). So, I calculate the test statistic, and I compare to the $\chi^2(15)$ distribution:

```
test1 <- 2 * (ll_I - ll_noI)
pval <- pchisq(test1, df = 15, lower.tail = FALSE)
sprintf("Test statistic: %.4f, p-value: %.4f", test1, pval)
```

```
## [1] "Test statistic: 18.4138, p-value: 0.2415"
```

Note that $\mathbb{P}(X^2 > \text{test}) = .2415 > .05$, so I fail to reject H_0 ; it seems as though the interaction terms are not jointly statistically significant.

c. Fitting Model without Interactions

i. Specify model, showing coefficients

I transform the age and car variables back into numeric variables and fit the no interactions model again:

```
claimsData2 <- claimsData
claimsData2$age <- as.numeric(claimsData2$age)
claimsData2$car <- as.numeric(claimsData2$car)
Inter2 <- glm(y ~ age + car + district + offset(log(n)),
              family = poisson(link = 'log'),
              data = claimsData2)
Inter2$coefficients
```

```
## (Intercept)      age      car  district1
## -1.8525316 -0.1767400  0.1977690  0.2186464
```

ii. Calculating Goodness of Fit statistic

I use the formula $X^2 = \sum_{i=1}^N \frac{(o_i - e_i)^2}{e_i}$:

```
expect <- exp(predict.glm(Inter2))
observe <- claimsData2$y
test2 = sum((observe - expect)^2/expect)
sprintf("Goodness of Fit Statistic: %.4f", test2)
```

```
## [1] "Goodness of Fit Statistic: 23.4976"
```

iii. Calculating deviance statistic

I use the formula $D = 2 \sum_{i=1}^N o_i * \log\left(\frac{o_i}{e_i}\right)$:

```
test3 <- 2 * sum(observe * log((observe + 0.000001)/expect))
sprintf("Deviance Statistic: %.4f", test3)
```

```
## [1] "Deviance Statistic: 24.6854"
```

2. Product binomial Distribution, Log-Linear

a. Proving simplification of algorithm

For arbitrary category i , let the distribution be modeled by Bernoulli(θ_{1i}) where θ_{1i} is the probability of success. Then, since $y_{.i} = n_i$ is fixed, given independent trials, I deduce that the random variable Z_i representing the number of successes in category i , is the sum of independent Bernoulli's - call them X_{ij} - so:

$$Z_i = \sum_{j=1}^{n_i} X_{ij} \sim \text{Binomial}(n_i, \theta_{1i}) = \text{Binomial}(n_i, \pi_i)$$

I assume each individual of the K categories is independent as to their distribution of success and failure. Letting $z_k = y_{1k}$, I deduce:

$$f(z_1, \dots, z_K | n_1, \dots, n_K) = \prod_{k=1}^K f_{Z_k}(z_k | n_k) = \prod_{k=1}^K \binom{n_k}{z_k} \pi_k^{z_k} (1 - \pi_k)^{n_k - z_k}$$

b. Proving log-linear equivalent to logistic

Note that, as $Z_k \sim \text{Binomial}(n_k, \pi_k)$, I deduce that $\mathbb{E}[Z_k] = n_k \pi_k$. Therefore:

$$\mathbb{E}[n_k - Z_k] = n_k - n_k \pi_k = n_k (1 - \pi_k)$$

From this, and the information given in the problem:

$$\begin{aligned} \log(\mathbb{E}[Z_k]) &= x_{1k}^T \beta \rightarrow \log(n_k \pi_k) = x_{1k}^T \beta \\ \log(\mathbb{E}[n_k - Z_k]) &= x_{2k}^T \beta \rightarrow \log(n_k (1 - \pi_k)) = x_{2k}^T \beta \end{aligned}$$

From this:

$$\begin{aligned} x_k^T \beta &= (x_{1k}^T - x_{2k}^T) \beta \\ &= x_{1k}^T \beta - x_{2k}^T \beta \\ &= \log(n_k \pi_k) - \log(n_k (1 - \pi_k)) \\ &= \log\left(\frac{n_k \pi_k}{n_k (1 - \pi_k)}\right) \\ &= \log\left(\frac{\pi_k}{1 - \pi_k}\right) \end{aligned}$$

c. Fitting Logistic, log-linear fits

First, I generate and fit the three logit models:

```
medData <- read_excel("Table 9.7 Ulcer and aspirin use.xls", skip = 2,
  sheet = "Sheet1", range = "A3:D11", .name_repair = "unique_quiet")
medDataLogit <- medData
medDataLogit$ulcer<-factor(medDataLogit$ulcer,levels = c("duodenal","gastric"),
  labels = c(0,1))
medDataLogit$`case-control` <- factor(medDataLogit$`case-control`,
  levels = c("control", "case"), labels = c(0,1))
medDataLogit$aspirin<- factor(medDataLogit$aspirin,
  levels = c("non-user", "user"), labels = c(0,1))
medLogitmodel <- glm(`case-control` ~ ulcer + aspirin,
  data = medDataLogit, weights = medDataLogit$frequency, family = 'binomial')
probs <- medLogitmodel$fitted.values[c(5,8,3,4)]
names(probs) <- c("Duo, noA", "Duo, A", "gas, noA", "gas, A")
probs
```

```
## Duo, noA Duo, A gas, noA gas, A
## 0.4464141 0.7166099 0.4204531 0.6946529
```

Now, I fit the log-linear model:

```
medData <- medData[,c("case-control", "ulcer", "aspirin", "frequency")]
medTab <- xtabs(frequency ~ `case-control` + aspirin + ulcer,
  data = medData)
logLin <- loglin(medTab, margin = list(c(1,2), c(1,3), c(2,3)), fit = TRUE)
```

```
## 4 iterations: deviation 0.01042226
```

```
logLin$fit
```

```
## , , ulcer = duodenal
##
##          aspirin
## case-control non-user user
## case      45.534222 11.465498
## control 56.465778  4.534502
##
## , , ulcer = gastric
##
##          aspirin
## case-control non-user user
## case      42.466498 21.533915
## control 58.533502  9.466085
```

I now check the fitted values for the probabilities:

```
probsLogLin<-c(logLin$fit[1]/sum(logLin$fit[1:2]),
  logLin$fit[3]/sum(logLin$fit[3:4]), logLin$fit[5]/sum(logLin$fit[5:6]),
  logLin$fit[7]/sum(logLin$fit[7:8]))
names(probsLogLin) <- c("Duo, noA", "Duo, A", "gas, noA", "gas, A")
probsLogLin
```

```
## Duo, noA Duo, A gas, noA gas, A
## 0.4464139 0.7165936 0.4204604 0.6946424
```

Both models produce similar results for ulcer probability based on type and aspirin use.

3. Calculating estimated beta, log-linear

I define the following variables:

$$x_{1i} = \begin{cases} 0 & \text{if Male} \\ 1 & \text{if Female} \end{cases} \quad \text{and} \quad x_{2i} = \begin{cases} 0 & \text{if Q} \\ 1 & \text{if R} \end{cases}$$

Therefore, using the predicted $\hat{\mu}$ values, I formulate 4 equations and 4 unknowns (implementing the values for x_1 and x_2 for each combination of categories):

$$\begin{aligned} \log(148) &= \hat{\beta}_0 \\ \log(446) &= \hat{\beta}_0 + \hat{\beta}_1 \\ \log(545) &= \hat{\beta}_0 + \hat{\beta}_2 \\ \log(4024) &= \hat{\beta}_0 + \hat{\beta}_1 + \hat{\beta}_2 + \hat{\beta}_3 \end{aligned}$$

Therefore:

$$\begin{aligned} \hat{\beta}_0 &= \log(148) \\ \rightarrow \hat{\beta}_1 &= \log(446) - \hat{\beta}_0 = \log(446) - \log(148) = \log\left(\frac{446}{148}\right) \\ \rightarrow \hat{\beta}_2 &= \log(545) - \hat{\beta}_0 = \log(545) - \log(148) = \log\left(\frac{545}{148}\right) \end{aligned}$$

Therefore, $\hat{\beta}_3 = \log(4024) - \hat{\beta}_0 - \hat{\beta}_1 - \hat{\beta}_2$, the estimated coefficient for the interaction term, can be calculated:

$$\hat{\beta}_3 = \log(4024) - \log(148) - \log\left(\frac{446}{148}\right) - \log\left(\frac{545}{148}\right) = \log\left(\frac{4024}{\frac{446*545}{148}}\right) = \mathbf{0.3982}$$

4. Chi-squared Goodness of Fit

a. Calculate the sample mean.

Let c_i be the number of policies with i claims. I calculate:

$$\bar{Y} = \frac{c_0(0) + c_1(1) + c_2(2) + c_3(3)}{c_0 + c_1 + c_2 + c_3} = \frac{450 + 80(2) + 20(3)}{2600} = \mathbf{0.2577}$$

b. Calculate the chi-square statistic

I first derive the maximum likelihood parameter prediction:

$$\begin{aligned} L(\lambda; \vec{y}) &= \prod_{i=1}^n f(y_i; \lambda) = \frac{e^{-n\lambda} * \lambda^{\sum_{i=1}^n y_i}}{\prod_{i=1}^n y_i!} \\ \rightarrow \ell(\lambda; \vec{y}) &= -n\lambda + \ln(\lambda) \sum_{i=1}^n y_i - \sum_{i=1}^n \ln(y_i!) \\ \rightarrow \ell'(\lambda; \vec{y}) &= -n + \frac{\sum_{i=1}^n y_i}{\lambda} = 0 \\ \rightarrow \frac{\sum_{i=1}^n y_i}{\lambda} &= n \\ \rightarrow \hat{\lambda}_{\text{M.L.E}} &= \frac{\sum_{i=1}^n y_i}{n} = \bar{Y} \end{aligned}$$

So $\hat{\lambda}_{\text{M.L.E}} = 0.2577$. Given the category breakdowns and the total number of policies, I calculate the expected number of claims:

$$\begin{aligned} \mathbb{E}[\text{number of 0 claim policies}] &= \mathbb{P}(y_i = 0) * n = e^{-.2577} * 2600 = 2009.366 \\ \mathbb{E}[\text{number of 1 claim policies}] &= \mathbb{P}(y_i = 1) * 2600 = .2577 * e^{-.2577} * 2600 = 517.798 \\ \mathbb{E}[\text{2 claim policies}] &= \mathbb{P}(y_i = 2) * 2600 = \frac{e^{-.2577} * .2577^2}{2} * 2600 = 66.716 \\ \mathbb{E}[\text{3 claim policies}] &= \mathbb{P}(y_i = 3) * 2600 = \frac{e^{-.2577} * .2577^3}{6} * 2600 = 5.731 \\ \mathbb{E}[\text{4+ claim policies}] &= \mathbb{P}(y_i > 3) * 2600 = \left(1 - \sum_{j=0}^3 \mathbb{P}(y_i = j)\right) * 2600 = 0.389 \end{aligned}$$

From these expected counts, and the observed counts given, I calculate the chi^2 statistic (letting k iterate over the categories):

$$\begin{aligned} X^2 &= \sum_{k=0}^4 \frac{(o_i - e_i)^2}{e_i} \\ &= \frac{(2050 - 2009.366)^2}{2009.366} + \frac{(450 - 517.798)^2}{517.798} + \frac{(80 - 66.716)^2}{66.716} + \frac{(20 - 5.731)^2}{5.731} + \frac{(0 - 0.389)^2}{0.389} \\ &= \mathbf{48.296} \end{aligned}$$

5. Predictions given fitted GLM

a. Calculate the predicted claim size

Note that, with a dispersion parameter of $\alpha = 1$:

$$\text{Claim Size} \sim \text{Exponential}(\hat{\theta}_{\text{M.L.E}})$$

Note that $\mathbb{E}[\text{Claim Size}] = \hat{\theta}_{\text{M.L.E}}$. Therefore, the model is as follows:

$$\begin{aligned} \ln(\hat{\theta}_{\text{M.L.E}}) &= \beta_0 + \beta_1 x_{z1,i} + \beta_2 x_{z2,i} + \beta_3 x_{z3,i} + \beta_4 x_{z5,i} \\ &\quad + \beta_5 x_{\text{convert},i} + \beta_6 x_{\text{coupe},i} + \beta_7 x_{\text{truck},i} + \beta_8 x_{\text{MV},i} + \beta_9 x_{\text{SW},i} + \beta_{10} x_{\text{U},i} \\ &\quad + \beta_{11} I_{\text{age} < 30,i} + \beta_{12} I_{\text{age} > 50,i} \end{aligned}$$

Therefore, for an observation of zone 3, with a truck and an age of 55:

$$\begin{aligned} \mathbb{E}[\text{Claim Size} | \text{Zone 3, Truck, 55}] &= \hat{\theta}_{\text{M.L.E}} \\ &= e^{\beta_0 + \beta_3 + \beta_7 + \beta_{12}} \\ &= e^{2.1 + 1.336 + 1.406 + 1.8} \\ &= \mathbf{766.627} \end{aligned}$$

b. Calculate Variance of claim size

For an exponential distribution, or a gamma distribution with an $\alpha = 1$, the dispersion parameter in the G.L.M is equal to $\lambda = 1$. Therefore, as the predicted claim severity is equal to $\hat{\theta}$ and, since, for $Y \sim \text{Exponential}(\theta)$, $\text{Var}[Y] = \theta^2$:

$$\text{Var}[\hat{\theta} | \text{Zone 4, Sedan, 35}] = 1 * \text{Var}[Y | \text{Zone 4, Sedan, 35}] = \hat{\theta}^2 | \text{Zone 4, Sedan, 35}$$

I first calculate the expected value:

$$\begin{aligned} \mathbb{E}[\text{Claim Size} | \text{Zone 4, Sedan, 35}] &= \hat{\theta}_{\text{M.L.E}} \\ &= e^{\beta_0} \\ &= e^{2.1} \\ &= 8.166 \end{aligned}$$

Therefore:

$$\text{Var}[Y | \text{Zone 4, Sedan, 35}] = 8.166^2 = \mathbf{66.686}$$